

Convolutional Neural Networks

FRC - EVL

2021

```
library(keras)
```

The Dogs vs. Cats dataset that you'll use isn't packaged with Keras. It was made available by Kaggle as part of a computer-vision competition in late 2013, back when convnets weren't mainstream. You can download the original dataset from www.kaggle.com/c/dogs-vs-cats/data. The pictures are medium-resolution color JPEGs.

Unsurprisingly, the cats-versus-dogs Kaggle competition in 2013 was won by entrants who used convnets. The best entries achieved up to 95% accuracy. In this example, you'll get fairly close to this accuracy (in the next section), even though you'll be training your models on less than 10% of the data that was available to the competitors.

This dataset contains 25,000 images of dogs and cats (12,500 from each class) and is 543 MB (compressed). After downloading and uncompressing it, you'll create a new dataset containing three subsets: a training set with 1,000 samples of each class, a validation set with 500 samples of each class, and a test set with 500 samples of each class.

Keras includes a number of image processing helper tools. In particular, it includes the *image_data_generator()* function, which can automatically turn image files on disk into batches of pre-processed tensors.

```
base_dir<-"~/Docencia/Curs_2020_2021_2S/UB/MESI0/SLT/Bloc2/CNN_lab/cats_and_dogs/"


train_dir<-file.path(base_dir,"train",fsep ="/")
validation_dir<-file.path(base_dir,"validation",fsep ="/")
test_dir<-file.path(base_dir,"test",fsep ="/")

train_cats_dir<-file.path(train_dir,"cats128",fsep ="/")
train_dogs_dir<-file.path(train_dir,"dogs128",fsep ="/")

validation_cats_dir<-file.path(validation_dir,"cats128",fsep ="/")
validation_dogs_dir<-file.path(validation_dir,"dogs128",fsep ="/")

test_cats_dir<-file.path(test_dir,"cats128",fsep ="/")
test_dogs_dir<-file.path(test_dir,"dogs128",fsep ="/")
```

CNN definition



```
model <- keras_model_sequential() %>%
  layer_conv_2d(filters = 32, kernel_size = c(3, 3), activation = "relu",
    input_shape = c(128, 128, 3)) %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 64, kernel_size = c(3, 3), activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 128, kernel_size = c(3, 3), activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 128, kernel_size = c(3, 3), activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_flatten() %>%
  layer_dense(units = 512, activation = "relu") %>%
  layer_dense(units = 1, activation = "sigmoid")
```

```
summary(model)
```

```
## Model: "sequential"
## -----
## Layer (type)                Output Shape                Param #
## -----
## conv2d_3 (Conv2D)           (None, 126, 126, 32)        896
## -----
## max_pooling2d_3 (MaxPooling2D) (None, 63, 63, 32)          0
## -----
## conv2d_2 (Conv2D)           (None, 61, 61, 64)          18496
## -----
## max_pooling2d_2 (MaxPooling2D) (None, 30, 30, 64)          0
## -----
## conv2d_1 (Conv2D)           (None, 28, 28, 128)         73856
## -----
## max_pooling2d_1 (MaxPooling2D) (None, 14, 14, 128)          0
## -----
## conv2d (Conv2D)             (None, 12, 12, 128)         147584
## -----
```

```
## max_pooling2d (MaxPooling2D)          (None, 6, 6, 128)          0
## -----
## flatten (Flatten)                    (None, 4608)          0
## -----
## dense_1 (Dense)                      (None, 512)          2359808
## -----
## dense (Dense)                        (None, 1)            513
## =====
## Total params: 2,601,153
## Trainable params: 2,601,153
## Non-trainable params: 0
## -----
```

```
# image_data_generator Generate batches of image data with real-time data augmentation. The data will b
train_datagen <- image_data_generator(rescale = 1/255) #
validation_datagen <- image_data_generator(rescale = 1/255) #

#flow_images_from_directory Generates batches of data from images in a directory (with optional augment
train_generator <- flow_images_from_directory(
  train_dir,
  train_datagen,
  target_size = c(128, 128),
  batch_size = 20,
  class_mode = "binary"
)
validation_generator <- flow_images_from_directory(
  validation_dir,
  validation_datagen,
  target_size = c(128, 128),
  batch_size = 20,
  class_mode = "binary"
)

batch <- generator_next(train_generator)
str(batch) # assigns image labels through the number of subdirectories (cat and dogs)

## List of 2
## $ : num [1:20, 1:128, 1:128, 1:3] 1 0.0392 0.5804 0.1725 0.4784 ...
## $ : num [1:20(1d)] 0 0 1 0 1 1 1 0 1 0 ...
```

Let's look at the output of one of these generators: it yields batches of 128×128 RGB images (shape (20, 128, 128, 3)) and binary labels (shape (20)). There are 20 samples in each batch (the batch size). Note that the generator yields these batches indefinitely: it loops endlessly over the images in the target folder.

```
model %>% compile(
  loss = "binary_crossentropy",
  optimizer = optimizer_rmsprop(lr = 1e-4),
  metrics = c("accuracy")
)
```

Let's fit the model to the data using the generator. You do so using the `fit_generator` function, the equivalent of `fit` for data generators like this one. It expects as its first argument a generator that will yield batches of inputs and targets indefinitely, like this one does. Because the data is being generated endlessly, the generator needs to know how many samples to draw from the generator before declaring an epoch over.

This is the role of the `steps_per_epoch` argument: after having drawn `steps_per_epoch` batches from the


generator—that is, after having run for `steps_per_epoch` gradient descent steps—the fitting process will go to the next epoch. In this case, batches are 20-samples large, so it will take 100 batches until you see your target of 2,000 samples. When using `fit_generator`, you can pass a `validation_data` argument, much as with the `fit` function. It's important to note that this argument is allowed to be a data generator, but it could also be a list of arrays. If you pass a generator as `validation_data`, then this generator is expected to yield batches of validation data endlessly; thus you should also specify the `validation_steps` argument, which tells the process how many batches to draw from the validation generator for evaluation.

```
history <- model %>% fit_generator(  
  train_generator,  
  steps_per_epoch = 100, #100  
  epochs = 5, # 20  
  validation_data = validation_generator,  
  validation_steps = 50 #50  
)
```

```
model %>% save_model_hdf5("cats_and_dogs_small_1.h5")
```

```
#model <- load_model_hdf5("cats_and_dogs_small_1.h5")
```

```
plot(history)
```



These plots are characteristic of overfitting. The training accuracy increases linearly over time, until it reaches nearly 100%, whereas the validation accuracy stalls at 71–75%. The validation loss reaches its minimum after only five epochs and then stalls, whereas the training loss keeps decreasing linearly until it reaches nearly 0. Because you have relatively few training samples (2,000), overfitting will be your number-one concern. You already know about a number of techniques that can help mitigate overfitting, such as dropout and weight decay (L2 regularization). We're now going to introduce a new one, specific to computer vision and used

almost universally when processing images with deep-learning models: data augmentation.

Using data augmentation

Overfitting is caused by having too few samples to learn from, rendering you unable to train a model that can generalize to new data. Given infinite data, your model would be exposed to every possible aspect of the data distribution at hand: you would never overfit. Data augmentation takes the approach of generating more training data from existing training samples, by augmenting the samples via a number of random transformations that yield believable-looking images. The goal is that at training time, your model will never see the exact same picture twice. This helps expose the model to more aspects of the data and generalize better. In Keras, this can be done by configuring a number of random transformations to be performed on the images read by an `image_data_generator`. Let's get started with an example.

These are just a few of the options available (for more, see the Keras documentation). Let's quickly go over this code:

- `rotation_range` is a value in degrees (0–180), a range within which to randomly rotate pictures.
- `width_shift` and `height_shift` are ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally.
- `shear_range` is for randomly applying shearing transformations.
- `zoom_range` is for randomly zooming inside pictures.
- `horizontal_flip` is for randomly flipping half the images horizontally—relevant when there are no assumptions of horizontal asymmetry (for example, real-world pictures).
- `fill_mode` is the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift.

```
datagen <- image_data_generator(  
  rescale = 1/255,  
  rotation_range = 40,  
  width_shift_range = 0.2,  
  height_shift_range = 0.2,  
  shear_range = 0.2,  
  zoom_range = 0.2,  
  horizontal_flip = TRUE,  
  fill_mode = "nearest",  
)  
  
fnames <- list.files(train_cats_dir, full.names = TRUE)  
img_path <- fnames[[3]]  
img <- image_load(img_path, target_size = c(128, 128))  
img_array <- image_to_array(img)  
img_array <- array_reshape(img_array, c(1, 128, 128, 3))  
augmentation_generator <- flow_images_from_data(  
  img_array,  
  generator = datagen,  
  batch_size = 1  
)  
  
op <- par(mfrow = c(2, 2), pty = "s", mar = c(1, 0, 1, 0))  
for (i in 1:4) {  
  batch <- generator_next(augmentation_generator) # Use to retrieve items from generators  
  plot(as.raster(batch[1,,,]))  
}
```



```
par(op)
```

```
datagen <- image_data_generator(  
  rescale = 1/255,  
  rotation_range = 40,  
  width_shift_range = 0.2,  
  height_shift_range = 0.2,  
  shear_range = 0.2,  
  zoom_range = 0.2,  
  horizontal_flip = TRUE,  
  fill_mode = "nearest"  
)  
  
train_generator <- flow_images_from_directory(  
  train_dir,  
  datagen,  
  target_size = c(128, 128),  
  batch_size = 20,  
  class_mode = "binary"  
)  
  
validation_datagen <- image_data_generator(rescale = 1/255)  
  
validation_generator <- flow_images_from_directory(  

```

```

validation_dir,
validation_datagen,
target_size = c(128, 128),
batch_size = 20,
class_mode = "binary"
)

batch <- generator_next(train_generator)
str(batch)

## List of 2
## $ : num [1:20, 1:128, 1:128, 1:3] 0.166 0.745 0.211 0.728 0.37 ...
## $ : num [1:20(1d)] 1 1 0 1 1 1 0 0 0 1 ...

model %>% compile(
  loss = "binary_crossentropy",
  optimizer = optimizer_rmsprop(lr = 1e-4),
  metrics = c("acc")
)


history <- model %>% fit_generator(
  train_generator,
  steps_per_epoch = 100, #100,
  epochs = 5, # 20
  validation_data = validation_generator,
  validation_steps = 50 #50
)

model %>% save_model_hdf5("cats_and_dogs_small_2.h5")

#model <- load_model_hdf5("cats_and_dogs_small_2.h5")

plot(history)

```

Using a pretrained convnet

A common and highly effective approach to deep learning on small image datasets is to use a pretrained network. A pretrained network is a saved network that was previously trained on a large dataset, typically on a large-scale image-classification task. If this original dataset is large enough and general enough, then the spatial-feature hierarchy learned by the pretrained network can effectively act as a generic model of the visual world, and hence its features can prove useful for many different computer-vision problems, even though these new problems may involve completely different classes than those of the original task. For instance, you might train a network on ImageNet (where classes are mostly animals and everyday objects) and then repurpose this trained network for something as remote as identifying furniture items in images. Such portability of learned features across different problems is a key advantage of deep learning compared to many older, shallow-learning approaches, and it makes deep learning very effective for small-data problems.

In this case, let's consider a large convnet trained on the ImageNet dataset (1.4 million labeled images and 1,000 different classes). ImageNet contains many animal classes, including different species of cats and dogs, and you can thus expect to perform well on the cats-versus-dogs classification problem. You'll use the VGG16 architecture, developed by Karen Simonyan and Andrew Zisserman in 2014; it's a simple and widely used convnet architecture for ImageNet.⁸ Although it's an older model, far from the current state of the art and somewhat heavier than many other recent models, we chose it because its architecture is similar to what you're already familiar with and is easy to understand without introducing any new concepts.

```
conv_base <- application_vgg16(
  weights = "imagenet",
  include_top = FALSE,
  input_shape = c(128, 128, 3)
)
```



```
conv_base
```

```
## Model
## Model: "vgg16"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_1 (InputLayer)        [(None, 128, 128, 3)]      0
## -----
## block1_conv1 (Conv2D)        (None, 128, 128, 64)       1792
## -----
## block1_conv2 (Conv2D)        (None, 128, 128, 64)       36928
## -----
## block1_pool (MaxPooling2D)   (None, 64, 64, 64)         0
## -----
## block2_conv1 (Conv2D)        (None, 64, 64, 128)        73856
## -----
## block2_conv2 (Conv2D)        (None, 64, 64, 128)        147584
## -----
## block2_pool (MaxPooling2D)   (None, 32, 32, 128)        0
## -----
## block3_conv1 (Conv2D)        (None, 32, 32, 256)        295168
## -----
## block3_conv2 (Conv2D)        (None, 32, 32, 256)        590080
## -----
## block3_conv3 (Conv2D)        (None, 32, 32, 256)        590080
## -----
## block3_pool (MaxPooling2D)   (None, 16, 16, 256)        0
## -----
## block4_conv1 (Conv2D)        (None, 16, 16, 512)        1180160
## -----
## block4_conv2 (Conv2D)        (None, 16, 16, 512)        2359808
## -----
## block4_conv3 (Conv2D)        (None, 16, 16, 512)        2359808
## -----
## block4_pool (MaxPooling2D)   (None, 8, 8, 512)          0
## -----
## block5_conv1 (Conv2D)        (None, 8, 8, 512)          2359808
## -----
## block5_conv2 (Conv2D)        (None, 8, 8, 512)          2359808
## -----
## block5_conv3 (Conv2D)        (None, 8, 8, 512)          2359808
## -----
## block5_pool (MaxPooling2D)   (None, 4, 4, 512)          0
## =====
## Total params: 14,714,688
## Trainable params: 14,714,688
## Non-trainable params: 0
## -----
```

```
model <- keras_model_sequential() %>%
  conv_base %>%
  layer_flatten() %>%
  layer_dense(units = 256, activation = "relu") %>%
```

```

layer_dense(units = 1, activation = "sigmoid")

summary(model)

## Model: "sequential_1"
## -----
## Layer (type)                Output Shape                Param #
## =====
## vgg16 (Model)                (None, 4, 4, 512)          14714688
## -----
## flatten_1 (Flatten)          (None, 8192)                0
## -----
## dense_3 (Dense)              (None, 256)                 2097408
## -----
## dense_2 (Dense)              (None, 1)                   257
## =====
## Total params: 16,812,353
## Trainable params: 16,812,353
## Non-trainable params: 0
## -----
cat("This is the number of trainable weights before freezing",
    "the conv base:", length(model$trainable_weights), "\n")

## This is the number of trainable weights before freezing the conv base: 30
freeze_weights(conv_base)    # Note that in order for these changes to take effect, you must recompile t

cat("This is the number of trainable weights after freezing",
    "the conv base:", length(model$trainable_weights), "\n")

## This is the number of trainable weights after freezing the conv base: 4
summary(model)

## Model: "sequential_1"
## -----
## Layer (type)                Output Shape                Param #
## =====
## vgg16 (Model)                (None, 4, 4, 512)          14714688
## -----
## flatten_1 (Flatten)          (None, 8192)                0
## -----
## dense_3 (Dense)              (None, 256)                 2097408
## -----
## dense_2 (Dense)              (None, 1)                   257
## =====
## Total params: 16,812,353
## Trainable params: 2,097,665
## Non-trainable params: 14,714,688
## -----
train_datagen = image_data_generator(
    rescale = 1/255,
    rotation_range = 40,

```

```

width_shift_range = 0.2,
height_shift_range = 0.2,
shear_range = 0.2,
zoom_range = 0.2,
horizontal_flip = TRUE,
fill_mode = "nearest"
)

train_generator <- flow_images_from_directory(
  train_dir,
  train_datagen,
  target_size = c(128, 128),
  batch_size = 20,
  class_mode = "binary"
)

validation_datagen <- image_data_generator(rescale = 1/255)

validation_generator <- flow_images_from_directory(
  validation_dir,
  validation_datagen,
  target_size = c(128, 128),
  batch_size = 20,
  class_mode = "binary"
)


model %>% compile(
  loss = "binary_crossentropy",
  optimizer = optimizer_rmsprop(lr = 2e-5),
  metrics = c("accuracy")
)

model %>% save_model_hdf5("cats_and_dogs_small_3.h5")

#model <- load_model_hdf5("cats_and_dogs_small_3.h5")

plot(history)

```



Fine-tuning

Fine-tuning consists of unfreezing a few of the top layers of a frozen model base used for feature extraction, and jointly training both the newly added part of the model (in this case, the fully connected classifier) and these top layers. This is called fine-tuning because it slightly adjusts the more abstract representations of the model being reused, in order to make them more relevant for the problem at hand.

```
unfreeze_weights(conv_base, from = "block3_conv1")
summary(model)
```

```
## Model: "sequential_1"
## -----
## Layer (type)                Output Shape          Param #
## =====
## vgg16 (Model)                (None, 4, 4, 512)     14714688
## -----
## flatten_1 (Flatten)          (None, 8192)           0
## -----
## dense_3 (Dense)              (None, 256)            2097408
## -----
## dense_2 (Dense)              (None, 1)              257
## =====
## Total params: 16,812,353
## Trainable params: 16,552,193
## Non-trainable params: 260,160
## -----
```

```

#model %>% compile(
#  loss = "binary_crossentropy",
#  optimizer = optimizer_rmsprop(lr = 2e-5),
#  metrics = c("accuracy")
#)

#history <- model %>% fit_generator(
#  train_generator,
#  steps_per_epoch = 100,
#  epochs = 30,
#  validation_data = validation_generator,
#  validation_steps = 50
#)

```

Feature extraction

```

datagen <- image_data_generator(rescale = 1/255)
batch_size <- 20

extract_features <- function(directory, sample_count) {
  features <- array(0, dim = c(sample_count, 4, 4, 512))
  labels <- array(0, dim = c(sample_count))
  generator <- flow_images_from_directory(
    directory = directory,
    generator = datagen,
    target_size = c(128, 128),
    batch_size = batch_size,
    class_mode = "binary"
  )
  i <- 0
  while(TRUE) {
    batch <- generator_next(generator)
    inputs_batch <- batch[[1]]
    labels_batch <- batch[[2]]
    features_batch <- conv_base %>% predict(inputs_batch)
    index_range <- ((i * batch_size)+1):((i + 1) * batch_size)
    features[index_range,,] <- features_batch
    labels[index_range] <- labels_batch
    i <- i + 1
    if (i * batch_size >= sample_count)
      break
  }
  list(
    features = features,
    labels = labels
  )
}

train <- extract_features(train_dir, 1000)
validation <- extract_features(validation_dir, 500)
test <- extract_features(test_dir, 500)

```

```
reshape_features <- function(features) {
  array_reshape(features, dim = c(nrow(features), 4 * 4 * 512))
}
train$features <- reshape_features(train$features)
dim(train$features)
```

```
## [1] 1000 8192
```

```
validation$features <- reshape_features(validation$features)
dim(validation$features)
```

```
## [1] 500 8192
```

```
test$features <- reshape_features(test$features)
dim(test$features)
```

```
## [1] 500 8192
```

```
model <- keras_model_sequential() %>%
  layer_dense(units = 256, activation = "relu",
    input_shape = 4 * 4 * 512) %>%
  layer_dropout(rate = 0.5) %>%
  layer_dense(units = 1, activation = "sigmoid")
```

```
model %>% compile(
  optimizer = optimizer_rmsprop(lr = 2e-5),
  loss = "binary_crossentropy",
  metrics = c("accuracy")
)
```

```
history <- model %>% fit(
  train$features, train$labels,
  epochs = 30,
  batch_size = 20,
  validation_data = list(validation$features, validation$labels)
)
```


```
summary(model)
```

```
## Model: "sequential_2"
```

```
## -----
## Layer (type)                Output Shape          Param #
## =====
## dense_5 (Dense)             (None, 256)           2097408
## -----
## dropout (Dropout)           (None, 256)           0
## -----
## dense_4 (Dense)             (None, 1)             257
## =====
## Total params: 2,097,665
## Trainable params: 2,097,665
## Non-trainable params: 0
## -----
```

```
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



Visualizing what convnets learn

Visualizing intermediate activations

```
model <- load_model_hdf5("~/Docencia/Curs_2020_2021_2S/UB/CursDLOmics/sessio2_CNN_lab/cats_and_dogs_small.h5")
model
```

```
## Model
## Model: "sequential"
## -----
## Layer (type)                Output Shape          Param #
## -----
## conv2d_3 (Conv2D)           (None, 126, 126, 32)  896
## -----
## max_pooling2d_3 (MaxPooling2D) (None, 63, 63, 32)    0
## -----
## conv2d_2 (Conv2D)           (None, 61, 61, 64)   18496
## -----
## max_pooling2d_2 (MaxPooling2D) (None, 30, 30, 64)    0
## -----
## conv2d_1 (Conv2D)           (None, 28, 28, 128)  73856
## -----
## max_pooling2d_1 (MaxPooling2D) (None, 14, 14, 128)   0
## -----
## conv2d (Conv2D)             (None, 12, 12, 128)  147584
```



```
## -----
## max_pooling2d (MaxPooling2D)          (None, 6, 6, 128)          0
## -----
## flatten (Flatten)                    (None, 4608)            0
## -----
## dense_1 (Dense)                      (None, 512)            2359808
## -----
## dense (Dense)                        (None, 1)              513
## =====
## Total params: 2,601,153
## Trainable params: 2,601,153
## Non-trainable params: 0
## -----
img<-image_load(paste0(test_cats_dir,"/", "cats1512.jpg"))
img_tensor<-image_to_array(img)
img_tensor<-array_reshape(img_tensor,c(1,128,128,3))
img_tensor<-img_tensor/255
dim(img_tensor)

## [1] 1 128 128 3
plot(as.raster(img_tensor[1,,]))
```



Extracts the outputs of the top eight layers and creates a model that will return these outputs, given the model input.

```

layer_outputs <- lapply(model$layers[1:8], function(layer) layer$output)
activation_model <- keras_model(inputs = model$input, outputs = layer_outputs)
activations <- activation_model %>% predict(img_tensor)

```

```

first_layer_activation <- activations[[1]]
dim(first_layer_activation)

```

```
## [1] 1 126 126 32
```


Compare with summary of the model

```


plot_channel <- function(channel) {
  rotate <- function(x) t(apply(x, 2, rev))
  image(rotate(channel), axes = FALSE, asp = 1,
  col = terrain.colors(12))
}

```

```
plot_channel(first_layer_activation[1,,1])
```




```
plot_channel(first_layer_activation[1,,5])
```



```
plot_channel(first_layer_activation[1,,15])
```




```
plot_channel(first_layer_activation[1,,,20])
```




```
# access other layers
layer_activation <- activations[[8]]
dim(layer_activation)

## [1] 1 6 6 128
plot_channel(layer_activation[1,,1])
```



```
plot_channel(layer_activation[1,,:,50])
```



```
plot_channel(layer_activation[1,,120])
```




Make predictions

```
#img<-image_load(paste0(test_cats_dir, "/", "cats1574.jpg"))  
img<-image_load(paste0(test_dogs_dir, "/", "dogs1574.jpg"))  
img_tensor<-image_to_array(img)  
img_tensor<-array_reshape(img_tensor, c(1, 128, 128, 3))  
img_tensor<-img_tensor/255  
dim(img_tensor)
```

```
## [1] 1 128 128 3
```

```
plot(as.raster(img_tensor[1,,]))
```



```
model %>% predict(img_tensor)
```

```
##           [,1]  
## [1,] 0.9696485
```