

Deep Neural Networks with Keras in R

A. Sanchez, F. Reverter and E. Vegas

Deep Learning with R

- We assume familiarity with:
 - Statistical learning basics
 - * Classification/Regression problems,
 - * Model building, Model evaluation
 - Deep Neural Networks (even more basic)
 - * Artificial neural networks, layers, activation function, gradient methods.
- And we focus on **how to build and use deep neural networks using R**

Outline

- Which software (and Hw) for Deep Learning
 - Python vs R
 - Tensorflow, Keras, Pytorch and many more
- Deep learning, vectorization and Tensors
 - Vectorization for efficient computation
 - Tensors for data representation and manipulation
- The Machine learning (DL?) Workflow
 - The general ML workflow
 - A Keras pipeline

Software for Deep Learning

1. [TensorFlow](#)
2. [PyTorch](#)
3. [Keras](#)
4. [MXNet](#)
5. [Caffe](#)
6. [Theano](#)
7. [Microsoft Cognitive Toolkit \(CNTK\)](#)

TensorFlow

- An open-source deep learning framework
- Developed by Google
- with a comprehensive ecosystem of tools and resources for building and deploying machine learning models.

[link to TensorFlow](#)



TensorFlow


Pytorch

- Open-source deep learning framework
- known for its dynamic computational graph feature
- and user-friendly interface
- developed by Facebook's AI Research lab.
- [link to PyTorch](#)



Keras

- High-level neural networks API
- written in Python
- runs on top of TensorFlow, CNTK, or Theano,
- emphasizes simplicity and ease of use.
- [link to Keras](#)




Keras

Deep learning with R

- As of 2023, common approach is
 - Tensorflow + Keras from within R
 - Using python in the background
- Multiple possible installations

- Possibly, the simplest is go to:
 - [TensorFlow for R site](#)



Keras

The Keras pipeline

- Training and using a model in keras is intended to be done through the usual steps of a ML workflow



A keras cheat sheet


Deep Learning with Keras : : CHEAT SHEET



Intro

Keras is a high-level neural networks API developed with a focus on enabling fast experimentation. It supports multiple back-ends, including TensorFlow, CNTK and Theano.

TensorFlow is a lower level mathematical library for building deep neural network architectures. The **keras** R package makes it easy to use Keras and TensorFlow in R.



<https://tensorflow.rstudio.com>

<https://www.manning.com/books/deep-learning-with-r-second-edition>

The "Hello, World!" of deep learning

INSTALLATION

The **keras** R package uses the Python **keras** library. You can install all the prerequisites directly from R. https://keras.rstudio.com/reference/install_keras.html

```
library(keras)
install_keras()
```

See ?install_keras for GPU instructions

This installs the required libraries in an Anaconda environment or virtual environment 'r-tensorflow'.

Working with keras models

DEFINE A MODEL

keras_model() Keras Model

keras_model_sequential() Keras Model composed of a linear stack of layers

multi_gpu_model() Replicates a model on different GPUs

COMPILE A MODEL

compile(object, optimizer, loss, metrics = NULL) Configure a Keras model for training

FIT A MODEL

fit(object, x = NULL, y = NULL, batch_size = NULL, epochs = 10, verbose = 1, callbacks = NULL, ...) Train a Keras model for a fixed number of epochs (iterations)

fit_generator() Fits the model on data yielded batch-by-batch by a generator

train_on_batch() **test_on_batch()** Single gradient update or model evaluation over one batch of samples

EVALUATE A MODEL

evaluate(object, x = NULL, y = NULL, batch_size = NULL) Evaluate a Keras model

evaluate_generator() Evaluates the model on a data generator

PREDICT

predict() Generate predictions from a Keras model

predict_proba() and **predict_classes()** Generates probability or class probability predictions for the input samples

predict_on_batch() Returns predictions for a single batch of samples

predict_generator() Generates predictions for the input samples from a data generator

OTHER MODEL OPERATIONS

summary() Print a summary of a Keras model

export_savedmodel() Export a saved model

get_layer() Retrieves a layer based on either its name (unique) or index

pop_layer() Remove the last layer in a model

save_model_hdf5(); load_model_hdf5() Save/Load models using HDF5 files

serialize_model(); unserialize_model() Serialize a model to an R object

clone_model() Clone a model instance

freeze_weights(); unfreeze_weights() Freeze and unfreeze weights

CORE LAYERS

layer_input() Input layer

layer_dense() Add a densely-connected NN layer to an output

layer_activation() Apply an activation function to an output

layer_dropout() Applies Dropout to the input

layer_reshape() Reshapes an output to a certain shape

layer_permute() Permute the dimensions of an input according to a given pattern

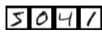
layer_repeat_vector() Repeats the input n times

layer_lambda(object, f) Wraps arbitrary expression as a layer

layer_activity_regularization() Layer that applies an update to the cost function based input activity

layer_masking() Masks a sequence by using a mask value to skip timesteps

layer_flatten() Flattens an input

input layer: use MNIST images 
mnist <- dataset_mnist()
x_train <- mnist\$train\$X; y_train <- mnist\$train\$y
x_test <- mnist\$test\$X; y_test <- mnist\$test\$y

reshape and rescale
x_train <- array_reshape(x_train, c(nrow(x_train), 784))
x_test <- array_reshape(x_test, c(nrow(x_test), 784))
x_train <- x_train / 255; x_test <- x_test / 255

y_train <- to_categorical(y_train, 10)
y_test <- to_categorical(y_test, 10)

defining the model and layers
model <- keras_model_sequential()
model %>%
 layer_dense(units = 256, activation = 'relu',
 input_shape = c(784)) %>%
 layer_dropout(rate = 0.4) %>%
 layer_dense(units = 128, activation = 'relu') %>%
 layer_dense(units = 10, activation = 'softmax')

compile (define loss and optimizer)

model %>% compile(
 loss = 'categorical_crossentropy',
 optimizer = optimizer_rmsprop(),
 metrics = c('accuracy')
)

train (fit)

model %>% fit(
 x_train, y_train,
 epochs = 30, batch_size = 128,
 validation_split = 0.2
)

model %>% evaluate(x_test, y_test)
model %>% predict_classes(x_test)



CC BY SA Posit Software, PBC • info@posit.co • posit.co • Learn more at tensorflow.rstudio.com • keras 2.11.1 • Updated: 2023-05

Available at [rstudio github repo](https://github.com/rstudio/keras)

Hello world of deep learning (1)

```
# load packages
library(keras)
# input layer: use MNIST images
mnist <- dataset_mnist()
x_train <- mnist$train$x; y_train <- mnist$train$y
x_test <- mnist$test$x; y_test <- mnist$test$y
```

Hello world of deep learning (2)

```
# reshape and rescale
x_train <- array_reshape(x_train, c(nrow(x_train), 784))
x_test <- array_reshape(x_test, c(nrow(x_test), 784))
x_train <- x_train / 255; x_test <- x_test / 255
y_train <- to_categorical(y_train, 10)
y_test <- to_categorical(y_test, 10)
```

Hello world of deep learning (3)

```
# defining the model and layers
model <- keras_model_sequential()
model %>%
  layer_dense(units = 256, activation = 'relu', input_shape = c(784)) %>%
  layer_dropout(rate = 0.4) %>%
  layer_dense(units = 128, activation = 'relu') %>%
  layer_dense(units = 10, activation = 'softmax')
```

Hello world of deep learning (4)

```
# compile (define loss and optimizer)
model %>% compile(
  loss = 'categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)
```


Hello world of deep learning (5)

```
# train (fit)
model %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 128,
  validation_split = 0.2
)
```

Hello world of deep learning (6)

```
model %>% evaluate(x_test, y_test)
model %>% predict_classes(x_test)
```

One must-do digression: Tensors

- Deep learning is filled with the word “tensor”,
 - Not to talk of *TensorFlow*
- What are Tensors any way?
 - R users: familiar with vectors (1-d arrays) and matrices (2-d arrays).
 - Tensors extend this concept to higher dimensions.
 - Can be seen as multi-dimensional arrays that generalize matrices.

Why tensors?


- Working with tensors has many benefits:
 - **Generalization:** Tensors generalize vectors and matrices to an arbitrary number of dimensions,
 - **Flexibility:** can hold a wide range of data types.
 - **Speed:** Use of tensors facilitates fast, parallel processing computations.

One and two dimensional tensors

Vectors: rank-1 tensors.




Matrices: rank-2 tensors.



Rank three tensors


- Arrays in layers.
- Typic use: Sequence data
 - time series, text
 - dim = (observations, seq steps, features)
- Examples

- 250 days of high, low, and current stock price for 390 minutes of trading in a day; $\text{dim} = \text{c}(250, 390, 3)$
- 1M tweets that can be 140 characters long and include 128 unique characters; $\text{dim} = \text{c}(1\text{M}, 140, 128)$



Rank four tensors

- Layers of groups of arrays
- Typic use: Image data
 - RGB channels
 - $\text{dim} = (\text{observations}, \text{height}, \text{width}, \text{color_depth})$
 - MNIST data could be seen as a 4D tensor where $\text{color_depth} = 1$



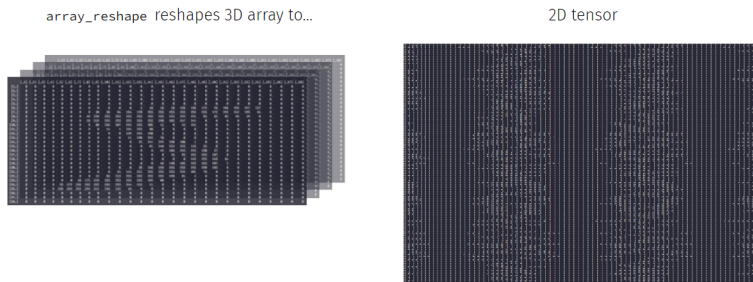
Rank five tensors

- Typic use: Video data
 - samples: 4 (each video is 1 minute long)
 - frames: 240 (4 frames/second)
 - width: 256 (pixels)
 - height: 144 (pixels)
 - channels: 3 (red, green, blue)
- Tensor shape (4, 240, 256, 144, 3)



One can always *reshape*

- Each DNN model has a given architecture which usually requires 2D/3D tensors.
- If data is not in the expected form it can be *reshaped*.



See [Deep learning with R](#) for more.

References and Resources

Workshops

- [Deep learning with R *Summer course*](#)
- [Deep learning with keras and Tensorflow in R \(Rstudio conf. 2020\)](#)

Books

- [Deep learning with R, 2nd edition. F. Chollet](#)

Documents

- [7 Best Deep Learning Frameworks To Watch Out For in 2022](#)